

# $Q$ -Learning for World Grid Navigation

**EE5904/ME5404 Part II: Project 2**

**Report due on 26 April 2019**

**Yáng YuánRùi**

yang\_yuanrui@u.nus.edu


11 April 2019

# Outline

- Project Description
- Recap
- Project Implementation
- Important Notes

# Project Description-Task

- The robot is to reach the goal state with maximum total reward of the trip



START


1	11	...	...	...	...	...	...	...	91
2	12	...	...	...	...	...	...	...	92
3	13							.	93
4	.							.	94
5	.							.	95
6	.							.	96
7	.							.	97
8	.							.	98
9	.							89	99
10	...	...	...	...	...	...	...	90	100



Illustration of a 10×10 world grid with start state and goal state.  
Index of each cell follows the Matlab column-wise convention for ease of programming

# Project Description: State Transition

- At a state, the robot can take 1 of the 4 actions
- The state transition is **deterministic**
- Assuming the state transition simulation is given by the **deterministic** state transition model
- You can actually use dynamic programming to find the optimal policy since the “model” is given
- In this project, you are only required to implement ***Q*-Learning**
- Some of the actions are not allowed, for the states moving out of the boundary



# Project Description: Reward Function

- Reward is given as a matrix “task1.mat” (known in Task 1)
- Reward Matrix:
  - Dimension:  $100 \times 4$
  - Each column represents an action (4 actions)
  - Each row represents a state (100 states)
- Prohibited transitions are marked by a reward of **-1**

# Recap

- Total Reward for an agent continuing its transition:

$$R_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \cdots = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1}$$

- $R_t$  determines present value of future rewards
- Rewards received  $k$  steps in the future is discounted by factor  $\gamma^{k-1}$
- Small  $\gamma$  forces agent to focus more on intermediate rewards from next few steps
- Large  $\gamma$  forces agent to take into account future rewards more strongly (agent becomes more farsighted)



# Recap

- ‘Worth’ actions at different states

$$Q^{\pi} : S \times A \rightarrow \mathcal{R}$$

- $Q^{\pi}(s, a) = E^{\pi}[R_t | s_t = s] \rightarrow R_t | s_t = s$

Deterministic Transition

- Expected return from taking action  $a$  at state  $s$  at time step  $t$  by following action  $\pi$
- Optimal policy is one that maximizes values of  $Q$ -functions overall all possible  $(s, a)$



# Recap: Optimal Policy

$$Q^\pi(s, a) = E^\pi [r_{t+1}] + E^\pi \left[ \gamma \sum_{k=0}^{\infty} \gamma^k r_{t+k+1} \mid s_t = s \right]$$

Slides 164-166

Values of  $Q$ -function are optimal if they are greater or equal to that of all other policies for all  $(s, a)$  pairs, i.e.,

$$Q^*(s, a) = \max_{\pi} Q^\pi(s, a)$$

## Greedy policy

At each  $s$ , select  $a$  that yields the largest value for the  $Q$ -function. When multiple choices are available, such  $a$  can be picked randomly

**Optimal policy:**  $\pi^*(s) \in \arg \max_a Q^*(s, a)$

Dynamic programming when state-transition model is given.



# Recap: Model-Free Value Iteration

When state transition model is **unknown**, the  $Q$ -function can be estimated via iterative update rule by using the reward received from observed state transition


$$Q_{k+1}(s_k, a_k) = Q_k(s_k, a_k) + \alpha_k \left( \underbrace{r_{k+1} + \gamma \max_{a'} Q_k(s_{k+1}, a')}_{\text{Estimate of } Q^*(s_k, a_k)} - Q_k(s_k, a_k) \right)$$

**Exploitation:** use **greedy** policy to select currently known best action

$$a_{k+1} = \max_{a'} Q_k(s_{k+1}, a')$$

**Exploration:** Try action other than current known best action

$$a_{k+1} \neq \max_{a'} Q_k(s_{k+1}, a')$$



**Exploitation:** Take  $a_4$

**Exploration:** Take  $a_1, a_2, a_3$

# Recap: $\epsilon$ -greedy exploration

## Initialization

Input: Discount factor  $\gamma$ ; exploration probability  $\epsilon_k$ ; learning rate  $\alpha_k$

- Initialize  $Q$ -function, e.g.,  $Q_0 \leftarrow 0$
- Determine the initial state  $s_0$
- For time step  $k$ , select action  $a_k$  according to:

## Select Action

$$a_k = \begin{cases} a \in \arg \max_{\hat{a}} Q_k(s_k, \hat{a}) & \text{Exploitation} \\ & \text{with probability } 1 - \epsilon_k \\ \text{an action uniformly randomly} & \\ \text{selected from all other actions} & \text{Exploration} \\ \text{available at state } s_k & \text{with probability } \epsilon_k \end{cases}$$

## Apply Action

## Update $Q$ -value

- Apply action  $a_k$ , receive reward  $r_{k+1}$ , then observe next state  $s_{k+1}$
- Update  $Q$ -function with:  
$$Q_{k+1}(s_k, a_k) = Q_k(s_k, a_k) + \alpha_k \left( r_{k+1} + \gamma \max_{a'} Q_k(s_{k+1}, a') - Q_k(s_k, a_k) \right)$$
- Set  $k = k + 1$  and repeat for-loop for the next time step

# Implementation: $Q$ -Learning

- Initialize
- For each trial
  - For each move
    - Select action
    - Apply action
    - Update  $Q$ -Value
- Extract Optimal policy

# Implementation: Parameter Setup

Input: Discount factor  $\gamma$ ; exploration probability  $\epsilon_k$ ; learning rate  $\alpha_k$

- Initialize  $Q$ -function, e.g.,  $Q_0 \leftarrow 0$
- Determine the initial state  $s_0$

TABLE I

PARAMETER VALUES AND PERFORMANCE OF  $Q$ -LEARNING

$\epsilon_k, \alpha_k$	No. of goal-reached runs		Execution time (sec.)	
	$\gamma = 0.5$	$\gamma = 0.9$	$\gamma = 0.5$	$\gamma = 0.9$
$\frac{1}{k}$	?	?	?	?
$\frac{100}{100+k}$	?	?	?	?
$\frac{1+\log(k)}{k}$	?	?	?	?
$\frac{1+5\log(k)}{k}$	?	?	?	?

$$\epsilon_k = \alpha_k$$

$k$  is time step

# Implementation: For each move

- Select action
- Apply action
- Update  $Q$ -Value

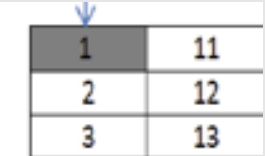
- For time step  $k$ , select action  $a_k$  according to:

$$a_k = \begin{cases} a \in \arg \max_{\hat{a}} Q_k(s_k, \hat{a}) & \text{with probability } 1 - \epsilon_k \\ \text{an action uniformly randomly} \\ \text{selected from all other actions} \\ \text{available at state } s_k & \text{with probability } \epsilon_k \end{cases}$$


- Apply action  $a_k$ , receive reward  $r_{k+1}$ , then observe next state  $s_{k+1}$
- Update  $Q$ -function with:  
$$Q_{k+1}(s_k, a_k) = Q_k(s_k, a_k) + \alpha_k \left( r_{k+1} + \gamma \max_{a'} Q_k(s_{k+1}, a') - Q_k(s_k, a_k) \right)$$
- Set  $k = k + 1$  and repeat for-loop for the next time step

# Implementation: For each move

- For the  $k^{\text{th}}$  movement, the machine is at state  $i$
- Choose the next movement based on  $\epsilon$ , (e.g. ,  $\epsilon = \frac{1}{k}$  )
- Assume the current optimal movement is to the left (i.e.  $a_k = 4$ )
  - Exploitation action is  $a_k = 4$ , with probability of  $1 - \epsilon = 1 - \frac{1}{k}$
  - Exploration action is  $a_k = 1, 2, 3$ , each with probability of  $\frac{\epsilon}{3} = \frac{1}{3k}$ 
    - What if  $i$  is at boundary?
    - Exploration action are uniformly selected from those possible explorative actions
- Assume  $a_k = 2$  is taken (i.e., move to the right)
  - $Q(i, 2) = Q(i, 2) + \alpha_k(\text{reward}(i, 2) + \gamma \max(Q(i + 10, :)) - Q(i, 2))$   
given from task1.mat



1	11
2	12
3	13



# Implementation: Termination Condition

*In theory:*

- **Trial termination condition:**

- In each trial, the robot starts at initial state ( $s = 1$ )
- It makes a series of transitions according to the algorithm for  $Q$ -learning with  $\epsilon$ -greedy exploration
- A trial ends when the robot reaches the goal state ( $s = 100$ )

- **Number of trials:**

- Repeat the process until the values of the  $Q$ -function converges to the optimal values

# Implementation: Termination Condition

*In this project:*

- **Trial termination condition:**
  - The robot reaches the goal state ( $s = 100$ ), *or*
  - $\alpha_k < 0.005$  (recommended)
    - You may also try other threshold
- **Number of trials:**
  - $Q$ -function converged to the optimal values, *or*
  - Number of trials  $\geq 3000$  (Try other options also)





# Implementation: Overview

- For each trial
  - For each move
    - Select action (perform exploitation or exploration)
    - Apply action
    - Update  $Q$ -value
$$Q_{next}(s, a) = Q_{crt}(s, a) + \alpha_k(\text{reward}(s, a) + \gamma \max(Q_{crt}(s', :)) - Q_{crt}(s, a))$$
  - Check trial termination condition
  - Check  $Q$ -value convergence / program termination condition
- Use the  $Q_{final}$  to extract the optimal path with greedy policy:
$$\pi^*(s) \in \arg \max_a Q^*(s, a)$$

# Implementation: Plot

- Show arrow
  - `plot(x, y, '^')%`, action 1
  - `plot(x, y, '>')%`, action 2
  - `plot(x, y, 'v')%`, action 3
  - `plot(x, y, '<')%`, action 4
- Show the grid world
  - plot lines to show the grid
- Subplot
  - easier viewing




For illustration only

# Important Notes: Task 1

TABLE I PARAMETER VALUES AND PERFORMANCE OF $Q$ -LEARNING				
$\epsilon_k, \alpha_k$	No. of goal-reached runs		Execution time (sec.)	
	$\gamma = 0.5$	$\gamma = 0.9$	$\gamma = 0.5$	$\gamma = 0.9$
$\frac{1}{k}$	?	?	?	?
$\frac{100}{100+k}$	?	?	?	?
$\frac{1+\log(k)}{k}$	?	?	?	?
$\frac{1+5\log(k)}{k}$	?	?	?	?

1. Matlab code are executable which generates reported results with the provided task1.mat
2. Complete Table I
3. Plot a  $10 \times 10$  grid showing the trajectory and the total reward
4. Provide necessary discussion based on your experimental results



# Important Notes: Task 2

- Choose a learning rate and discount rate wisely so that your robot can deal with the unknown reward.
- The code will be used to find the optimal policy using a reward function *not* provided to the students.
- “**qeval.mat**” will be used (as a replacement for the reward you have from task1.mat) to evaluation your RL program.

# Important Notes: Task 2

- You can assume that the unknown reward “**qevalreward**” are loaded in the workspace
- The state transition model, initial state, goal state, are the same as those in Task 1, but the reward function is different.
- The output of your program should be a column vector named “**qevalstates**” to store the trajectory
- Also plot a  $10 \times 10$  grid showing the trajectory and the total reward

# Important Notes: Task 1 & 2

- The program evaluation will be based on
  - Policy
  - Execution time
  - *Executable!*
- You can create your own reward matrix to test the effectiveness of your code
- The marking will be based on the **report** and the **program code**. Explain your choice of parameters clearly in your report.

# Important Notes

- Name your RL main script (for Task 2) as “**RL\_main**” for testing unknown reward.
- Please play around with the discount factor, the learning rate and the exploration probability.
- Use your **student number** as the folder name. Generate a non-password-protected **zipfile** of this folder and upload this zipfile to the IVLE.

**Report due on 26 April 2019**

Thank you!